

Econometrics in R's tidyverse

by Tyler Ransom, University of Oklahoma
[@tyleransom](#)

Basics of R

R can be thought of as a really fancy calculator

Packages:

- R comes with a lot of functionality out-of-the-box
- Other functionality requires the user to load packages
- One-time installation: `install.packages("tidyverse")`
- Each time you open R: `library(tidyverse)`

Commenting:

- Use `#` to make a comment
- This tells R to ignore that code
`# My name is Tyler`

Assignment operator:

- Use `<-` to store a calculation, e.g. `x <- 3` (" $x = 3$ ")

Pipe operator:

- Use `%>%` to "pipe" objects
`y <- mean(log(x))` becomes `y <- x %>% log %>% mean`
- `%<>%` pipes forward, then backwards
`x <- mean(log(x))` is same as `x %<>% log %>% mean`

Working with Data

R's fundamental data object is a **data frame**

Like spreadsheets, stores data in columns and rows

tidyverse uses **tibbles** (enhanced data frames)

```
df <- as_tibble(mtcars)
```

Reading in data

- Many functions for reading in different types of data
`df <- read_csv("myfile.csv")` (comma separated)
`df <- read_fwf("myfile.dat")` (fixed-width)
- More details: see [Data Importing Cheat Sheet](#)
- **haven** package: import foreign files (e.g. SAS, Stata, ...)

Accessing columns of data

- To reference a column in a tibble, use `$`
`df$mpg`
`mean(df$mpg)` will return sample avg of mpg variable

Ignore missing values

- Missing values are indicated by `NA`
- Some commands won't automatically ignore `NA` values

- For these cases, use `na.rm` option
`mean(df$mpg, na.rm=TRUE)`
`df$mpg %>% mean(na.rm=TRUE)` (equivalent)
- Otherwise, R would say the mean is `NA`

Removing columns and rows from a tibble

- To keep columns in a tibble, use `select()`
`df1 <- df %>% select(mpg, disp, hp, gear, carb)`
- To keep rows in a tibble, use `filter()`
`df1 %<>% filter(mpg >= 10)`
- To remove columns, put a minus in front
`df1 <- df %>% select(-mpg, -disp)`

Remove missing values from a tibble

- To remove **all** rows with *any* `NA` values, use `drop_na()`
`df1 <- df %>% drop_na()`
- Can also drop `NA`'s from particular columns:
`df1 <- df %>% drop_na(gear, carb)`

Creating new columns in a tibble

- To create a new column in a tibble, use `mutate()`
`df1 %<>% mutate(mpg_squared = mpg^2)`

Manipulating values of a variable

- To replace (i.e. recode) values of a variable:
`df %<>% mutate(gear = replace(gear, gear==4, 99))`
Changes all 4's in `gear` to be 99's
`gear==4` can be any other logical condition
- To specify a series of conditions, use `%in%`
`df %<>% mutate(hp = replace(hp, hp %in% c(110, 120), 99))`
Changes all 110's or 120's in `hp` to be 99's

Working with discrete variables

- Discrete variables often require special treatment
- In R, declare discrete variables as **factors**
`df %<>% mutate(gear = as.factor(gear))`

Other data manipulations

- See [Data Wrangling Cheat Sheet](#)

Getting to know your data

It's important to know what's in your data by

1. Looking at summary statistics
2. Performing cross-tabulations
3. Visualizing certain variables

Summary statistics

- Report quartiles, min/max, mean, and `#NA`'s:
`summary(df)`
- Report `N` (total non-missing), mean, SD, min/max:
`df %>% as.data.frame %>% stargazer(type="text")`

Cross-tabulations

- Report frequencies of a discrete variable:
`table(df$gear)`
- Average y by categories of a discrete x variable:
`df %>% group_by(gear) %>% summarize(m.mpg = mean(mpg))`

Visualization

- Often helpful to look at a histogram or line graph
- Histogram (continuous x):
`ggplot(df, aes(mpg)) + geom_histogram()`
- Histogram (factor x):
`ggplot(df, aes(x=gear)) + geom_bar()`
- Kernel density plot:
`ggplot(df, aes(mpg)) + geom_density()`
- Simple scatter plot with linear fit:
`ggplot(df, aes(disp, mpg)) + geom_point() + geom_smooth(method="lm")`
- More details: see [ggplot2 Cheat Sheet](#)

Regression modeling

Basic OLS regression

- Regression:

```
est <- lm(mpg ~ gear + hp, data=df)
```
- Examine regression output:

```
summary(est)  
tidy(est)  
stargazer(est,type="text")
```
- Other functional forms:

```
est <- lm(mpg ~ gear + I(gear^2), data=df)  
est <- lm(log(mpg) ~ gear + I(gear^2), data=df)
```
- Factor variables automatically get separate intercepts

t-statics and *F*-statistics

- *t*-stats, *p*-values reported in regression output
- *F*-test:

```
linearHypothesis(est,c("gear","hp"))  
tests  $H_0 : \beta_{gear} = 0, \beta_{hp} = 0$   
  
linearHypothesis(est,c("gear=5","hp=-1"))  
tests  $H_0 : \beta_{gear} = 5, \beta_{hp} = -1$ 
```
- Robust *F*-test (see next section):

```
linearHypothesis(est.rob,c("gear","hp"))
```

Robust standard errors (estimatr package)

- Correct for heteroskedasticity:

```
est.rob <- lm_robust(mpg ~ gear + hp, data=df)  
or  
stargazer(est,se=starpred(est.rob),type="text")
```
- Correct for serial correlation:

```
fixed.est <- est %>% coeftest(vcov=NeweyWest)  
stargazer(est,se=list(fixed.est[,2]),type="text")
```
- Correct for clustering:

```
est.clust <- lm_robust(mpg ~ gear + hp, data=df,  
clusters=df$carb)  
or  
stargazer(est,se=starpred(est.clust),type="text")
```

Instrumental Variables

- Let `drat` be the endogenous covariate
- Let `wt` be the instrument
- Let `qsec` and `gear` be exogenous covariates

```
est.iv <- ivreg(mpg ~ drat + qsec + gear |  
wt + qsec + gear, data=df)
```
- Instruments come after the `|` symbol
- Endogenous covariates come before the `|` symbol
- Exogenous covariates appear on both sides of the `|`
- First-stage regression:

```
est.1 <- lm(drat ~ wt + qsec + gear, data=df)  
df %<>% mutate(drat.hat = est.1$fitted.values)
```
- Second-stage regression:

```
est.2 <- lm(mpg ~ drat.hat + qsec + gear, data=df)
```
- Can also use `estimatr` for robust SEs:

```
est.ivr <- iv_robust(mpg ~ drat + qsec + gear |  
wt + qsec + gear, data=df)
```

Working with time series data

- Declare a time series data frame

```
df.ts <- zoo(df, order.by=df$year)
```
- Time series line plot:

```
ggplot(df.ts, aes(year, inf)) + geom_line()
```
- Simple AR(1) model:

```
est <- dynlm(inf ~ L(inf,1), data=df.ts)
```
- First-differences model:

```
est.diff <- dynlm(d(inf) ~ unem, data = df.ts)
```
- ADF test for unit root:

```
adf.test(df1.ts$inf, k=1)
```
- ARIMA model:

```
est.arima <- auto.arima(df.ts$inf)
```
- Plot *h*-period-ahead forecast intervals

```
autoplot(forecast(est.arima, h=2))
```
- Extended date and time functions available in
[lubridate package](#)

Working with panel data

- Report number of units and time periods

```
pdim(df)
```
- Pooled OLS model

```
est.pols <- plm(lwage ~ exper + I(exper^2) +  
year, data = df, index = c("id","year"),  
model = "pooling")
```
- Random effects model

```
est.re <- plm(lwage ~ exper + I(exper^2) +  
year, data = df, index = c("id","year"),  
model = "random")
```
- Fixed effects model

```
est.fe <- plm(lwage ~ exper + I(exper^2) +  
year, data = df, index = c("id","year"),  
model = "within")
```
- First differences model

```
est.fd <- plm(lwage ~ exper + I(exper^2) +  
year, data = df, index = c("id","year"),  
model = "fd")
```

Limited dependent variable models

Linear probability model (LPM):

- If *y* is a factor, format it as a numeric

```
est.lpm <- lm(as.numeric(y) ~ x1 + x2, data=df)
```

Logit and Probit:

In this case, *y* should be formatted as a factor

- Logit:

```
est.logit <- glm(y ~ x1 + x2,  
family=binomial(link="logit"),data=df)
```
- Probit:

```
est.probit <- glm(y ~ x1 + x2,  
family=binomial(link="probit"),data=df)
```

List of packages

The document requires the following packages:

tidyverse	car	sandwich	tseries
magrittr	estimatr	zoo	lubridate
stargazer	lmtest	dynlm	forecast
broom	clubSandwich	AER	plm

Layout by Winston Chang, <http://wch.github.io/latexsheet/>